

# CS4416 Database system project

Christoffer Näs, Jules Espinoux, Jeremy Nicolas, David Chuntishvili

December 1, 2024

## Contents

<b>1</b>	<b>Contribution</b>	<b>1</b>
<b>2</b>	<b>Platform</b>	<b>1</b>
<b>3</b>	<b>Modifications to schema</b>	<b>1</b>
<b>4</b>	<b>Entity Relationship Diagram</b>	<b>2</b>
<b>5</b>	<b>View and triggers</b>	<b>2</b>
5.1	View . . . . .	2
5.2	Triggers . . . . .	3
<b>6</b>	<b>Function and procedures</b>	<b>3</b>
6.1	total_nr_of_occupied_seats function . . . . .	3
6.2	add_song_album_link procedure . . . . .	3
<b>7</b>	<b>Index requirement for efficient execution</b>	<b>3</b>
7.1	Indexes for the triggers . . . . .	3
7.2	Indexes for the function . . . . .	3
7.3	Indexes for the procedure . . . . .	3

## 1 Contribution

All work was done together and distributed equally.

## 2 Platform

Platform used for the project was XAMPP on Windows 11.

## 3 Modifications to schema

This project aimed to improve the existing database to better handle the details of artists, albums, songs, concerts, and fans interactions. In the updated schema, each table has a unique primary key to make sure every record is distinct and can be quickly found. This helps keep the database accurate and organized.

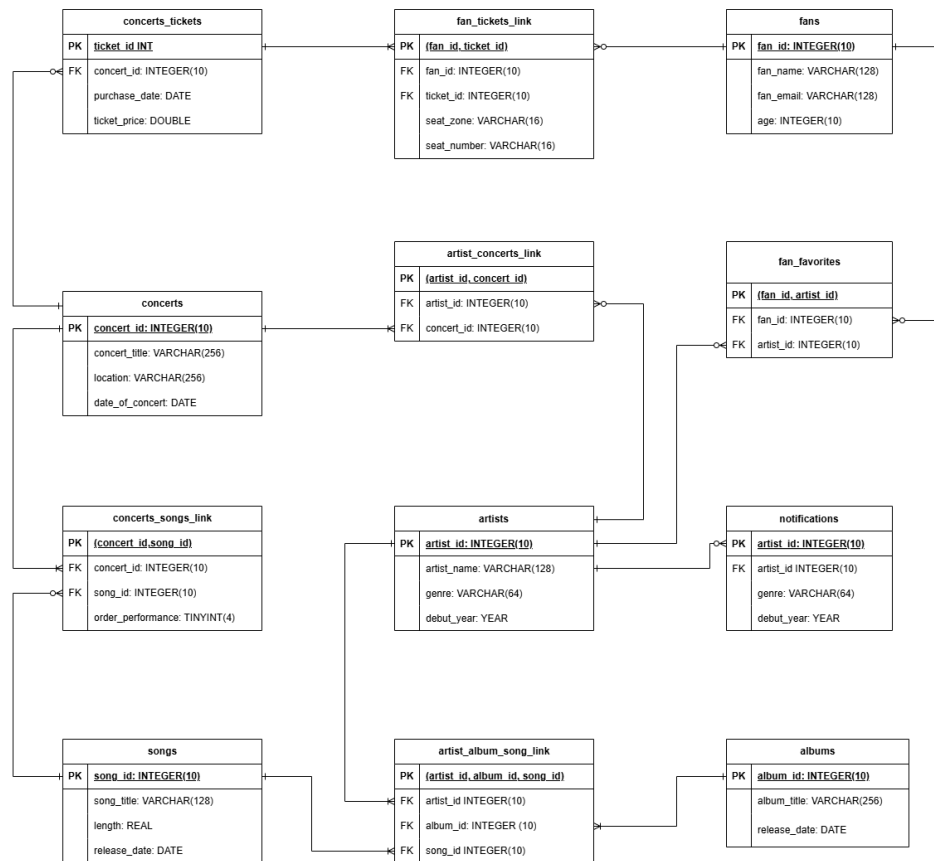
Originally, the schema had direct relationships, like linking artists straight to albums, which made it hard to show multiple artists working on the same album. To fix this, we added new tables like "artist\_album\_song\_link" that allow for many artists to be connected to many albums. This change makes the database more flexible and true to how music albums are actually made.

This table also let us show that multiple artists can work on one song and that many songs can be played at one concert. These updates help the database better reflect the teamwork and dynamics of music production and concerts.

Updating the fan data was another key improvement. Before, fan details were directly put into the

Together, these changes enhance the database's functionality and efficiency, making sure it can handle a wider range of real-life situations.

Our ERD shows how the data in our system is organized and connected. It maps out the key parts of the data, like entities, their details, and how they relate to each other. It's a clear plan for how the data fits together, without getting into how it's physically stored in the database.



The view `concerts_over_one_place_summary` summarizes concerts with more than one ticket sold, showing the concert title, total songs performed, total song duration, and tickets sold. Data comes from several tables: `concerts`, `concerts_songs.link`, `songs`, `concerts_tickets`, and `artist_concerts.link`. The

query uses JOINS to connect these tables by IDs, groups the data by concert title, and calculates totals for songs, duration, and tickets. It filters out concerts with one or fewer tickets sold and sorts results by tickets sold.

## 5.2 Triggers

The trigger `fan_deletion` makes sure so all the links connected to a fan is removed before a record in the fan table is removed making sure that all references to the fan is removed on deletion. The second trigger `notify_artist_on_new_concert` works in the following way, after signing a new artist to a concert and adding it to the database, an automatic notification record is created in the database that could be used to send out to the artist or to all of his fans.

## 6 Function and procedures

### 6.1 `total_nr_of_occupied_seats` function

This function returns the number of occupied seats for a given concert. To do this, the function takes all the tickets associated to a concert and count the number of fan (and then seat) associated with each ticket.

### 6.2 `add_song_album_link` procedure

This procedure checks whether a given song is associated with a given album. To do this, we do the assumption that a song stored in the database will always have an associated value in the table `artist_album_song_link`. Then we can retrieve the artist associated to the song and add the link if it's not already in the table.

## 7 Index requirement for efficient execution

Introducing indexes to tables and columns where data retrieval is performed can improve query execution time since data is localized faster. However introducing indexes on tables make modification of the database slower since all the indexes would need to be updated on modification.

### 7.1 Indexes for the triggers

For trigger `fan_deletion`, using indexes would make the execution of the trigger slower since it is modifying the database by performing deletion of rows.

For the trigger `notify_artist_on_new_concert`, using indexes for table `concerts` and column `date_of_concert` would optimize our query retrieval time since we are retrieving rows from this column in the trigger. However introducing indexes for the table `notifications` would slower our deletions on the table notifications and also weaken our trigger execution time since modifications on the database is performed on insertion of rows.

### 7.2 Indexes for the function

Introducing indexes for table `concerts_tickets` and column `ticket_id` would improve the function `total_nr_of_occupied_seats` execution time since it is retrieving data from this column.

### 7.3 Indexes for the procedure

To improve execution time for procedure `add_song_album_link`, indexes could be added to table `songs` and the column `release_date` since the columns is only used for data retrieval. For the other columns used the rows are also modified so there is a trade off to think about. The procedure is more read heavy so introducing indexes for column `release_date` in `albums`, column `song_id`, `album_id` in `artist_album_song_link` is probably worth it.